

JAGS example

Paul M

3/27/2020

Introduction to JAGS (“Just Another GIBBS Sampler”)

[Based on one of the examples at <http://www.johnmyleswhite.com/notebook/2010/08/20/using-jags-in-r-with-the-rjags-package/>]

First you must install jags on your computer! rjags won't run without it. e.g., in the terminal using homebrew: brew install jags" Then install the rjags package:

```
#install.packages("rjags")
library(rjags)
```

```
## Loading required package: coda
## Linked to JAGS 4.3.0
## Loaded modules: basemod,bugs
library(coda)
```

An example using normal rvs, where we are inferring the mean and variance:

Let's assume we have a set of rvs that we know are normally distributed, but we wish to infer their mean and variance. Generate a test set of data

```
N <- 1000
x <- rnorm(N, 0, 2)

write.table(x,
            file = 'example1.data',
            row.names = FALSE,
            col.names = FALSE)
```

Now we need to write a model specification in JAGS syntax. Put the model specification (below) in a file called example1.bug. The complete model looks like this:

```
model { for (i in 1:N) {
  x[i] ~ dnorm(mu, tau)
}
mu ~ dnorm(0, .0001)
tau <- pow(sigma, -2)
sigma ~ dunif(0, 100)
}
```

- The first line says that you are specifying a model.
- Then you set up the model for every single data point using a for loop. Here, we say that $x[i]$ is distributed normally with mean μ and precision τ . Note that rjags works with precision rather than variance! (Precision=1/variance) *Then we specify priors for μ and τ .
- μ is assumed to be distributed normally with mean 0 and standard deviation 100. This is an example of a non-informative prior.
- Then we specify τ as a deterministic function (hence the deterministic <- instead of the distributional ~) of σ , after raising σ to the -2 power. Then we say that σ has a uniform prior over the interval [0,100].

Now we invoke the model:

```
# Set up our model object in R
jags <- jags.model('example1.bug',  # specification file
                  data = list('x' = x,  # the data (must use same names as trhe model spec. file)
                              'N' = N),
                  n.chains = 4,  # how many parallel chains to run
                  n.adapt = 100  # we will use adaptive sampling, removing the first 100 iterations
                  )

## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 1000
##   Unobserved stochastic nodes: 2
##   Total graph size: 1009
##
## Initializing model

update(jags, 1000)  # run another 1000 iterations. The update function is used to add more iterations

samps <- jags.samples(jags,  # draw 1000 samples from the output for the requested variables
                     c('mu', 'tau'),
                     1000)

summary(samps)

##      Length Class   Mode
## mu   4000    marray numeric
## tau  4000    marray numeric
```

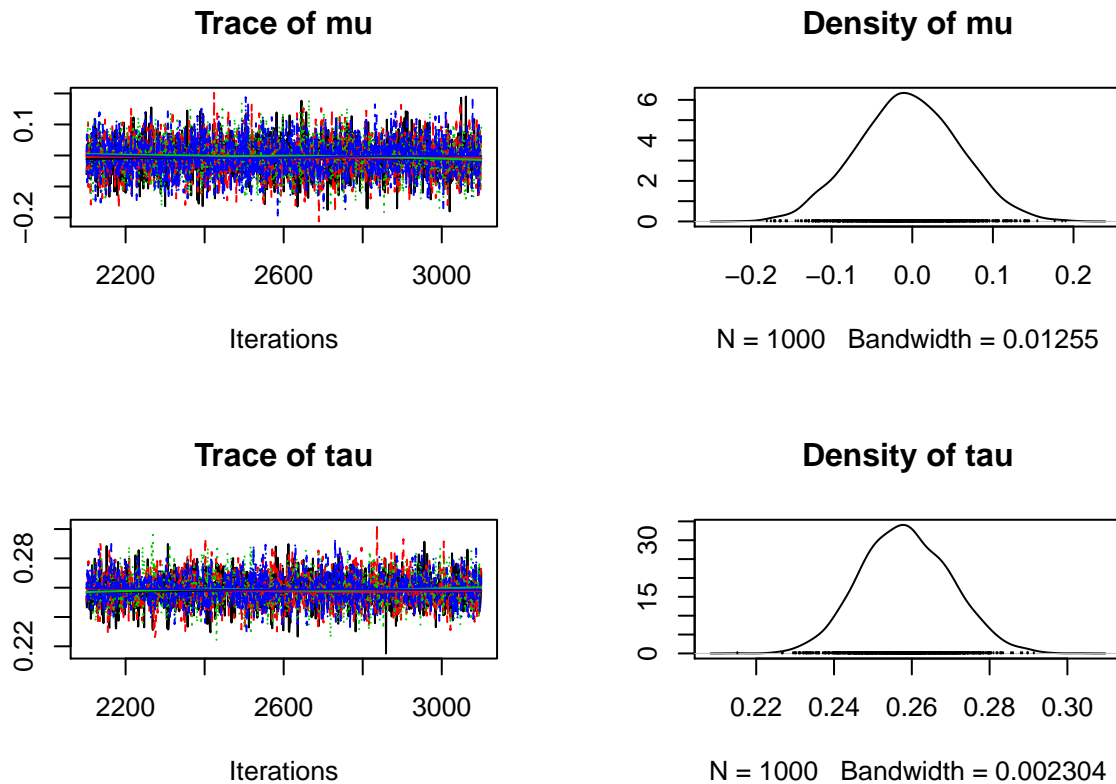
You can also set it for use of coda diagnostics for convergence (which also seems to provide a better summary of the output)

```
samps2 <- coda.samples( jags, c('mu','tau'), 1000 )
summary(samps2)

##
## Iterations = 2101:3100
## Thinning interval = 1
## Number of chains = 4
## Sample size per chain = 1000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
```

```
##           Mean      SD Naive SE Time-series SE
## mu  -0.005113 0.06220 0.0009834    0.0009886
## tau  0.258463 0.01142 0.0001805    0.0001949
##
## 2. Quantiles for each variable:
##
##      2.5%      25%      50%      75%  97.5%
## mu  -0.1248 -0.04706 -0.005097 0.03765 0.1167
## tau  0.2369  0.25051  0.258155 0.26632 0.2811
```

```
plot(samps2)
```



```
show(gelman.diag(samps2))
```

```
## Potential scale reduction factors:
##
##      Point est. Upper C.I.
## mu           1         1
## tau           1         1
##
## Multivariate psrf
##
## 1
```